

Algoteq Brown Bag series

Pilot Episode - Apache Ignite

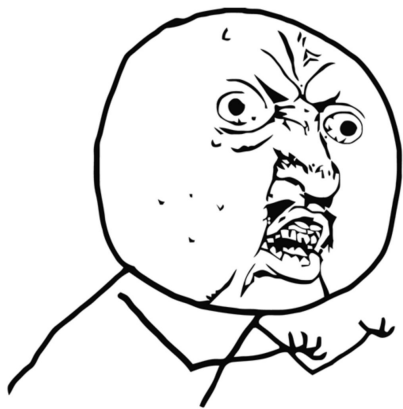
Ignite - what's that?

- A distributed database - can hold data on multiple physical machines
- Key-value pairs storage (Redis style), with optional SQL access
- Distributed transactions supported
- In-memory-centric, with optional persistence
- Written in Java



So... what?

So it's just a Java-written in-memory database... right?



Y waste our time then?!?!!!



Because...

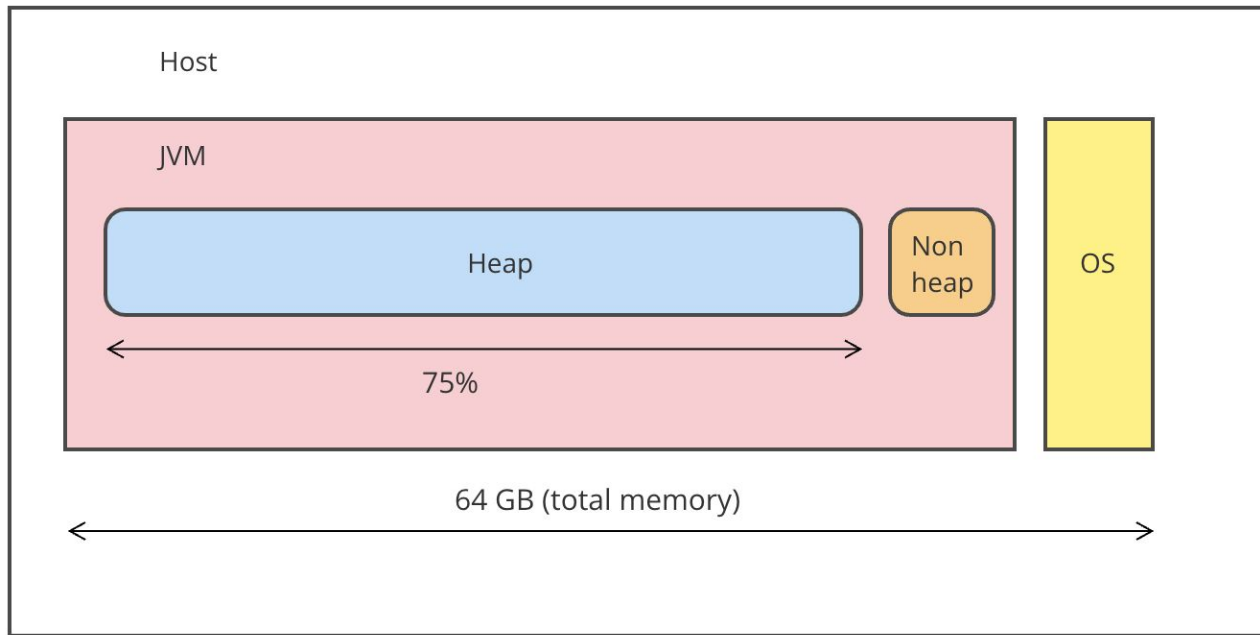
because OFF-HEAP

... because DATA COLOCATION & FAILOVER

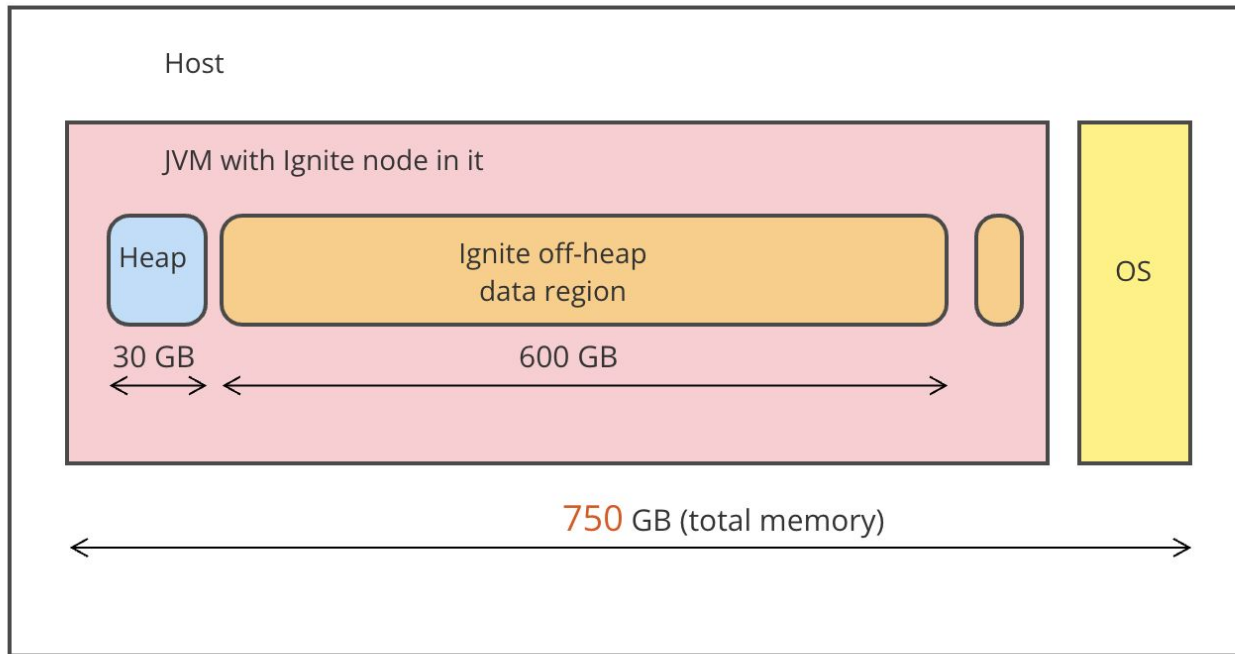
... because INTEGRATED SERVICES



Typical Java server



Single (!!!) PROD Ignite node



Show me the code

(code)

(code)(code)

(code)(code)(code)(code)

(code)(code)

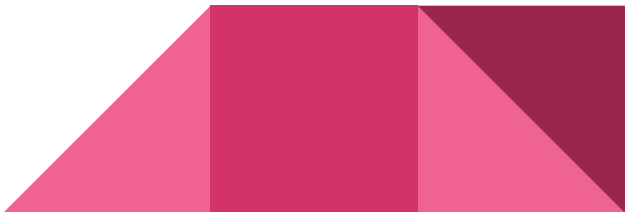
(code)(code)

(code)(code)(code)

(code)(code)(code)(code)

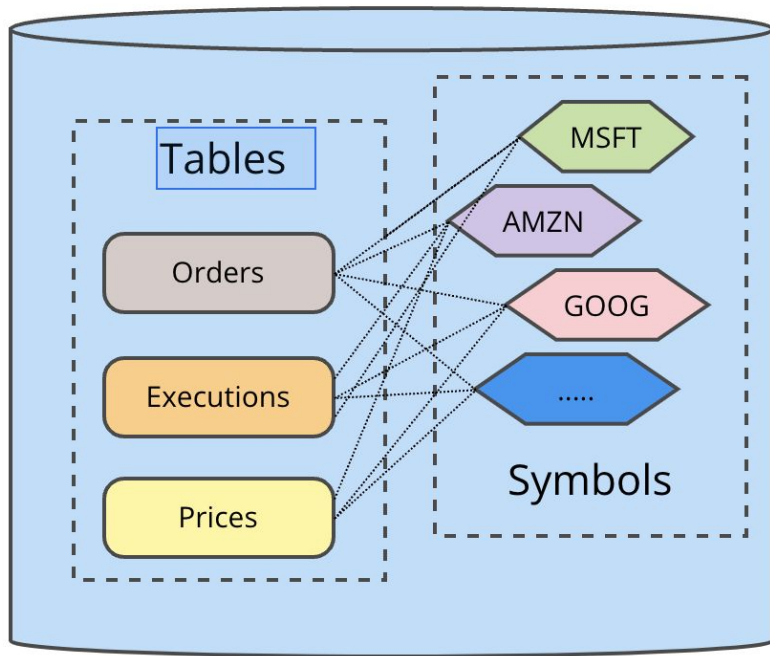
(code)(code)(code)(code)(code)

(code)(code)(code)(code)(code)(code)

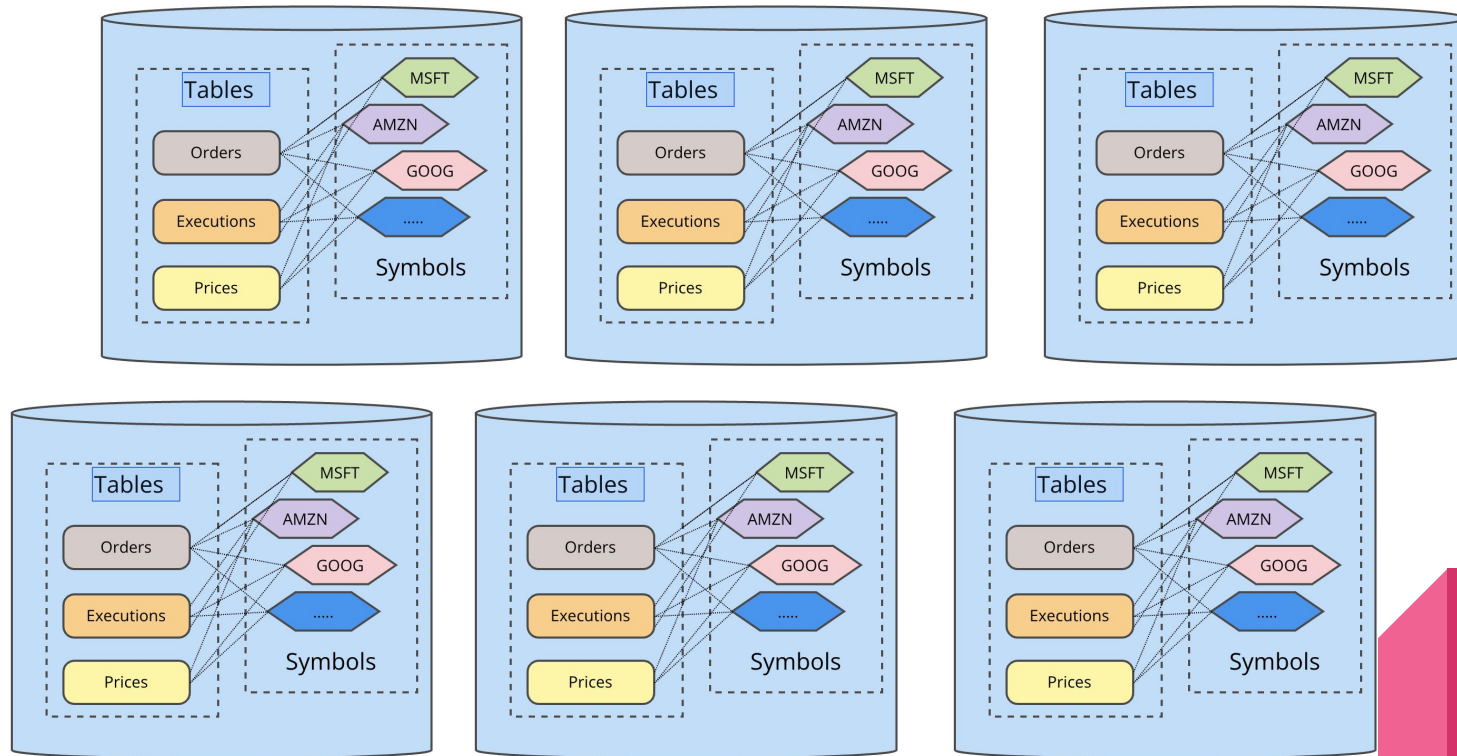


Data colocation - why

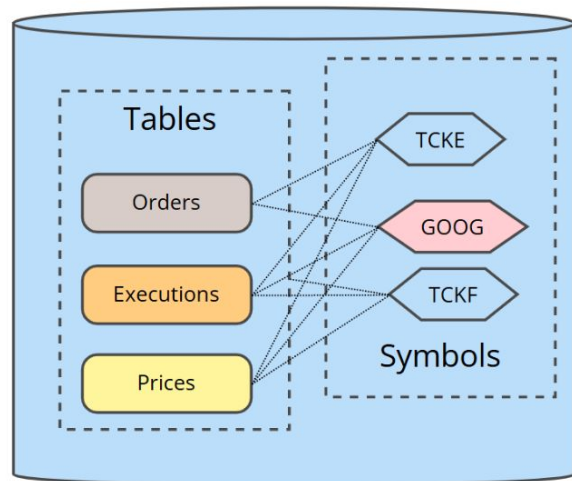
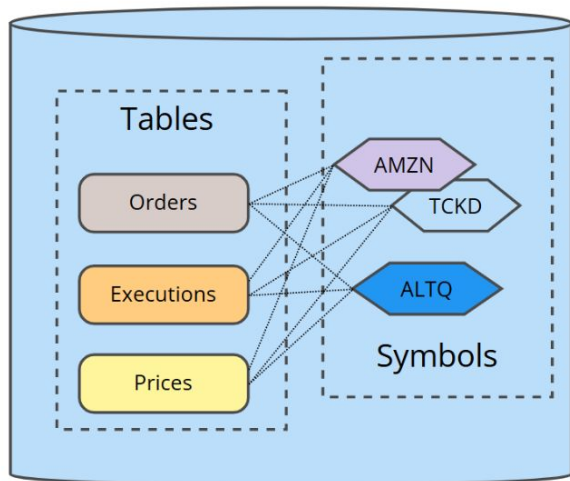
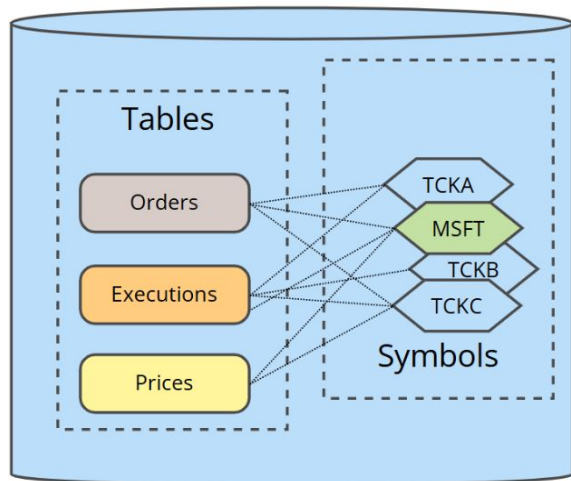
Consider millions of fast-ticking records here:



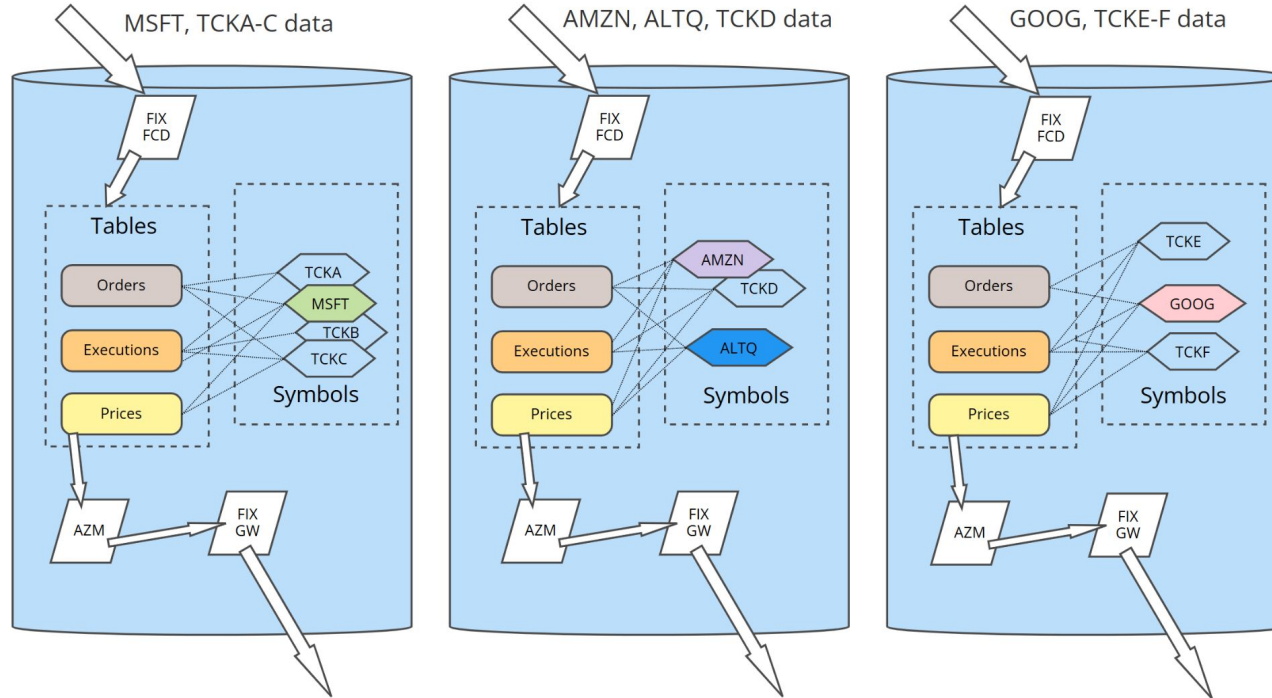
Data colocation - why



Data colocation - that's why (A)



Data colocation - that's why (B)



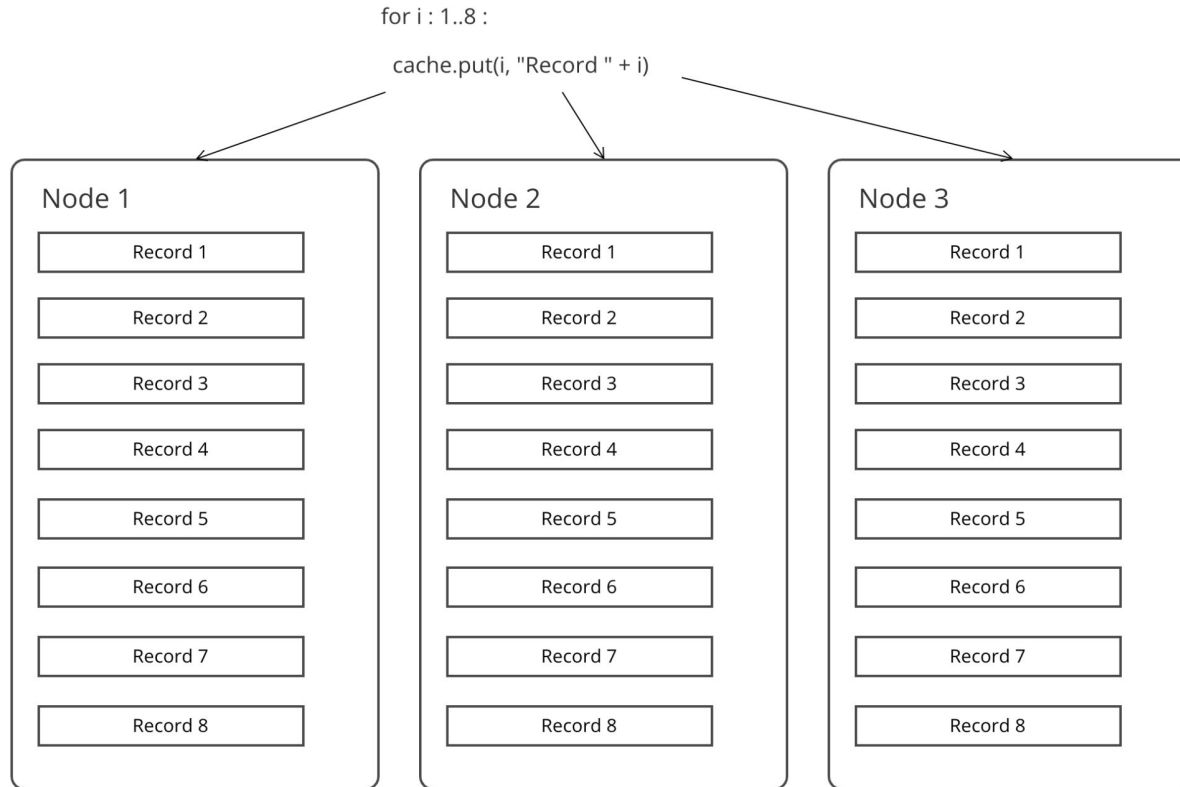
Colocation - how

Ignite caches can have one of two mode:

- REPLICATED
- - or -
- PARTITIONED

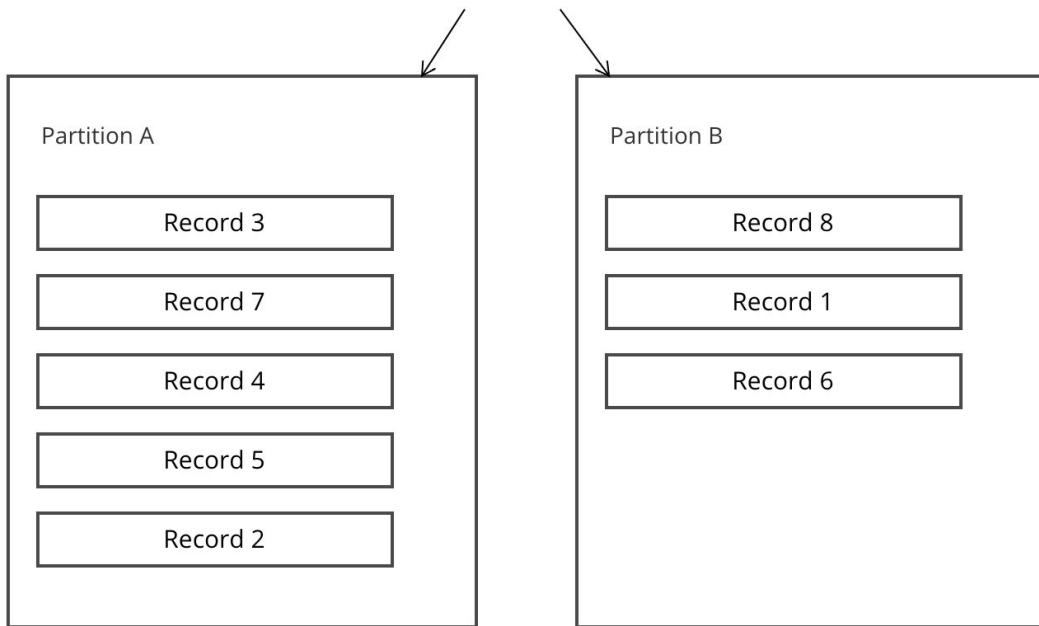


REPLICATED cache

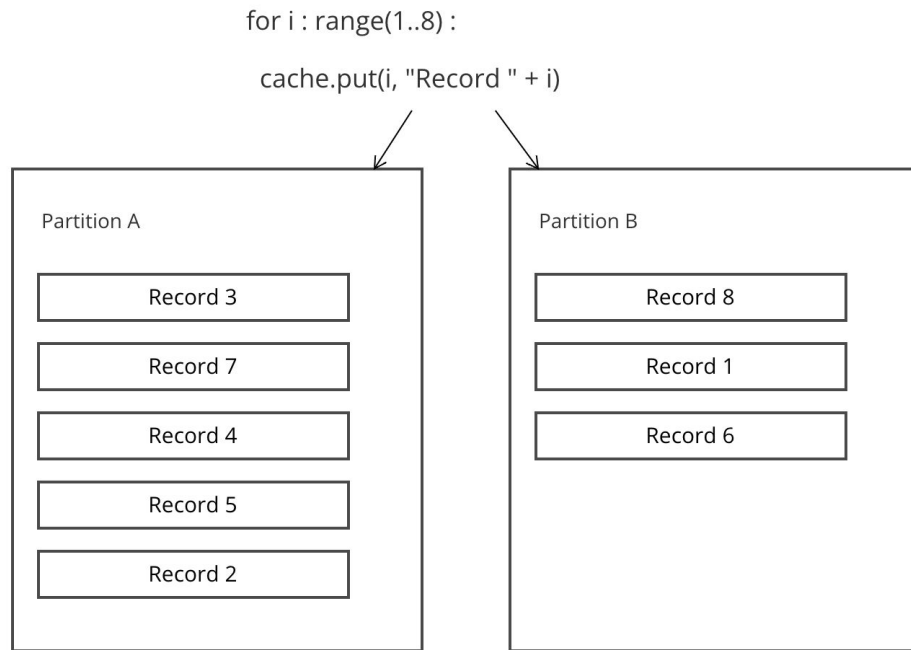


What is a partition?

```
for i : range(1..8):  
  cache.put(i, "Record " + i)
```

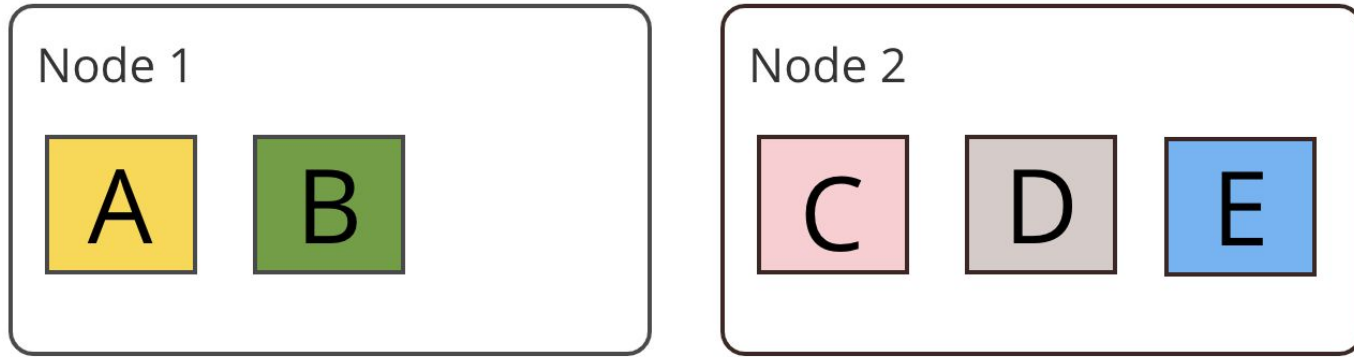


What is a partition?

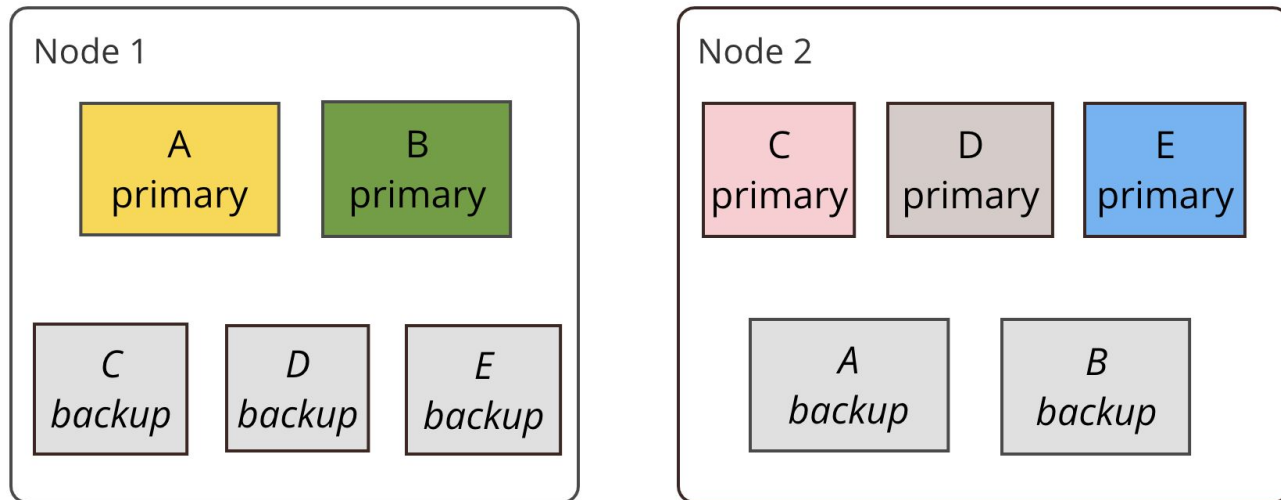


- Cache consists of the set of records.
- Partition is a logical subset of those records.
- Partition is a `_stable_` subset, once an ID is assigned to a partition, it never moves to another one.

Simplest case - partitions with no back-ups

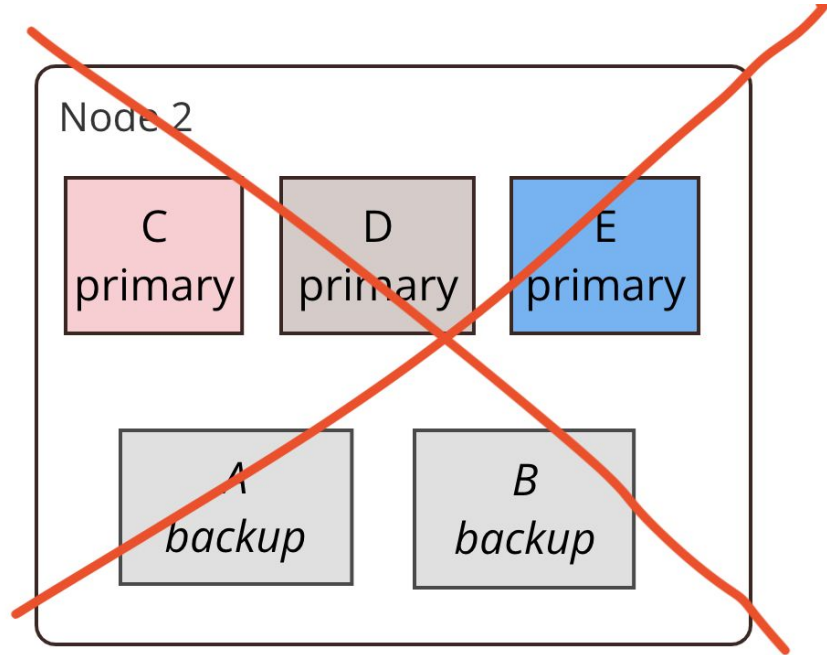
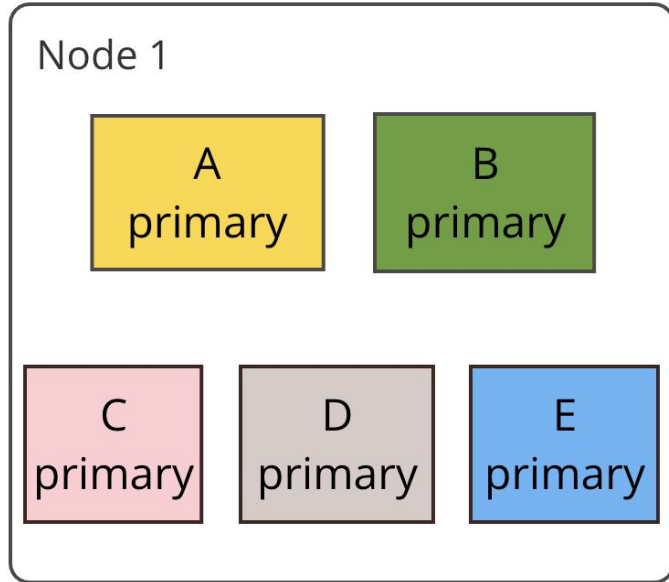


Partitions with backup



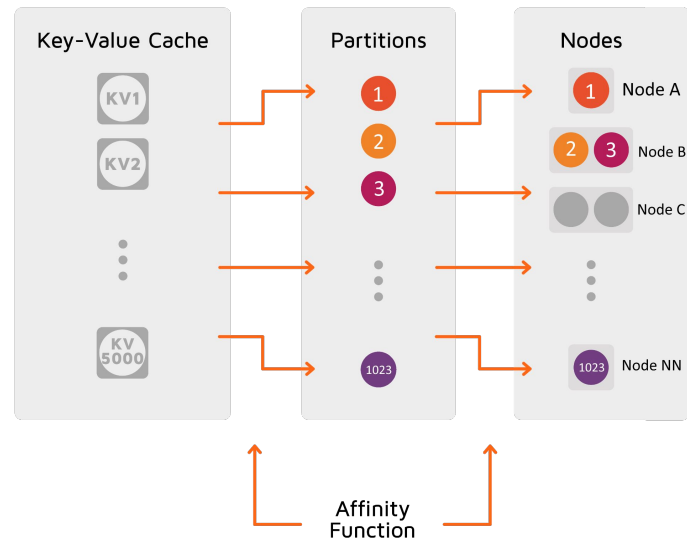
UBER RULE: a partition is never split between the nodes

Partitions with backup - fail-over after a node killed



Affinity function


- The affinity function determines the mapping between keys and partitions.
- Partition are identified by a number from 0 to 1023.
- The set of partitions is distributed between the server nodes available at the moment.
- When the number of nodes in the cluster changes, the partitions are re-distributed.
- The affinity function takes the affinity key as an argument.



Affinity key

```
8 public class DataKey {  
9     public long id;  
10  
11     @AffinityKeyMapped  
12     public final String affinityKey;  
13  
14     public DataKey(long id, String affinityKey) {  
15         this.id = id;  
16         this.affinityKey = affinityKey;  
17     }  
18
```

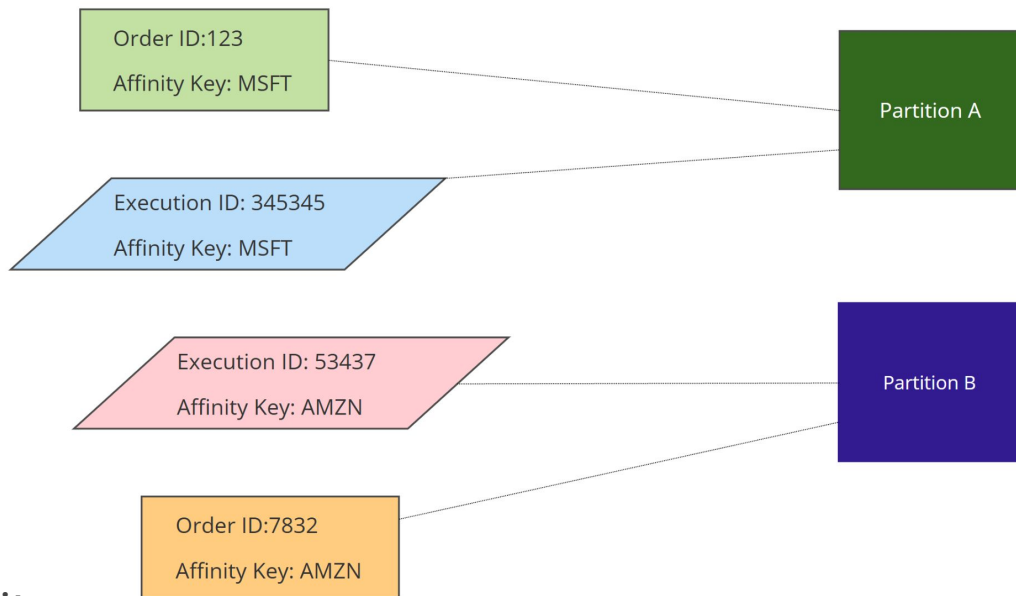
**declares
affinity key**



Cache with affinity key sample

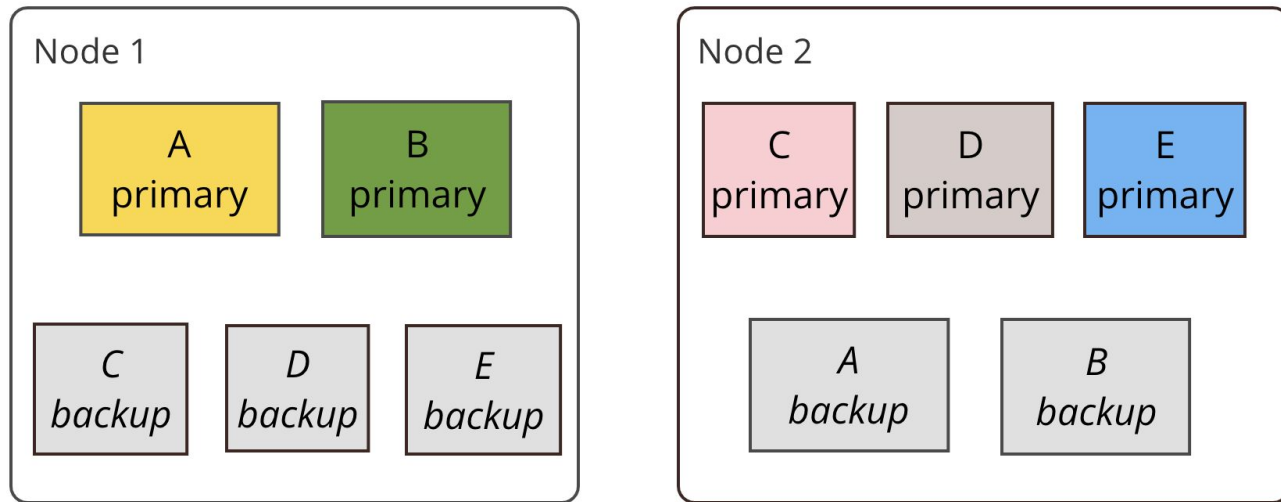
[illegible]

Data co-location



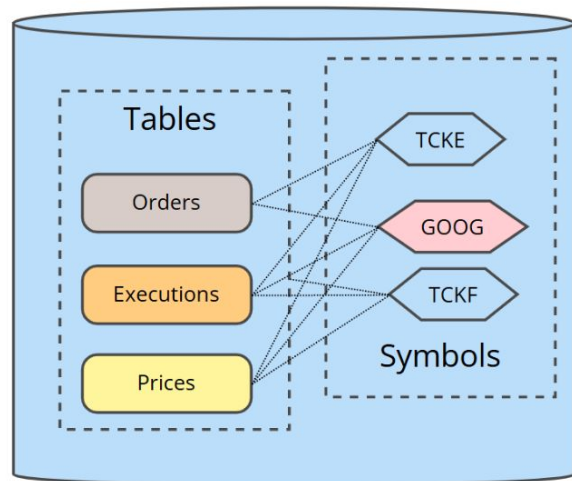
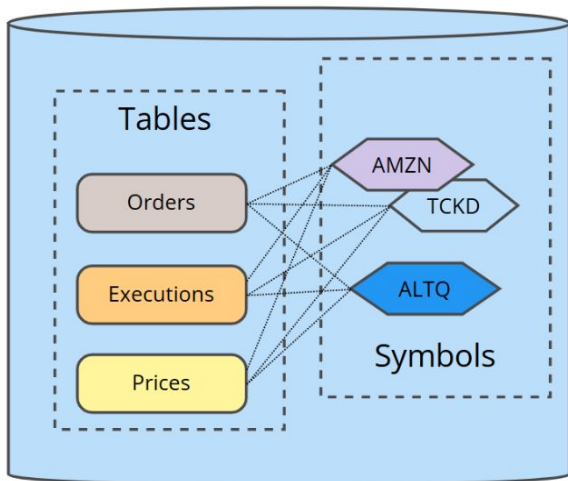
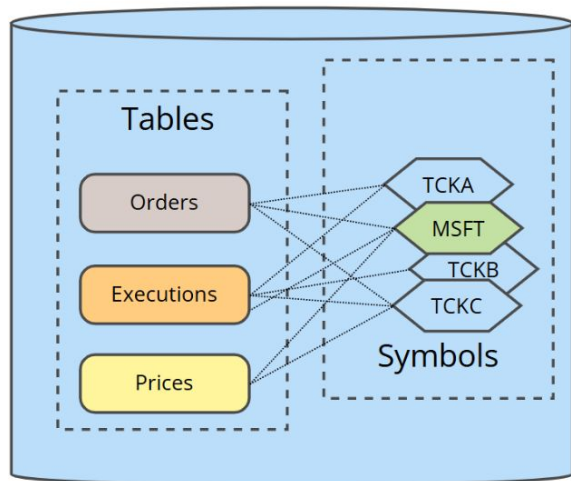
Affinity function looks at the affinity key only, NOT the cache type

Partitions with backup



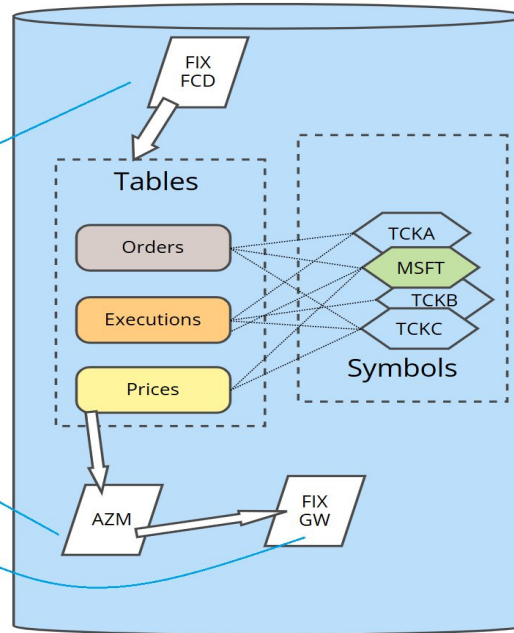
UBER RULE: a partition is never split between the nodes

Data is colocated now



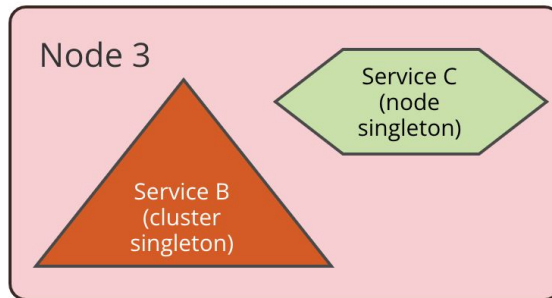
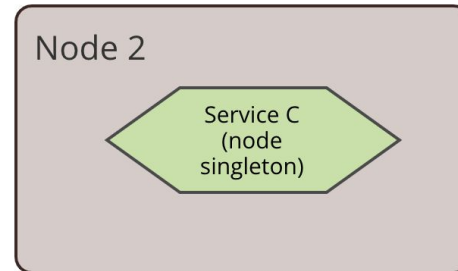
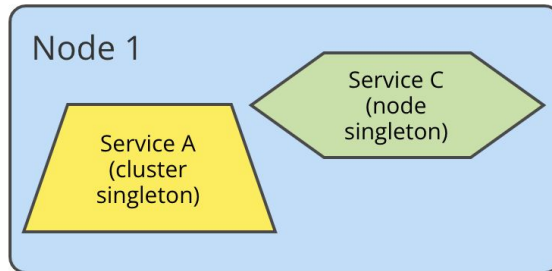
Services

Services



Many types of services-to-nodes-config

- Cluster singleton
- Node singleton
- Affinity services
- ...



Services rebalancing

